

Functional JDK 17

Jan Ypma

October 26, 2021

Contents

1	Introduction	2
1.1	About today's presentation	2
2	Recent new JDK features	3
2.1	Records (Java 16)	3
2.1.1	Demo	3
2.1.2	Summary	4
2.2	Sealed classes	4
2.2.1	Implemented in different source files	4
2.2.2	Implemented in the same source file	4
2.2.3	Can't extend outside of what's sealed in	4
2.3	"Pattern matching" for switch	5
2.3.1	Switch in Java 7	5
2.3.2	Java 12: Strings, expressions	5
2.3.3	Java 12: Arrows for expressions	5
2.3.4	Java 13: Yield instead of break	5
2.3.5	Java 17: Pattern matching objects, and guards	6
2.3.6	Exhaustiveness check	6
2.3.7	Case branch for null	6
2.3.8	Limitations	6
2.4	Shenandoah GC and ZGC	7
2.4.1	Traditional garbage collectors	7
2.4.2	ZGC and Shenandoah GC	7
2.4.3	ZGC	7
2.4.4	Shenandoah GC	7
2.5	macOS / AArch64 port	8
3	Practical reactive streams	8
3.1	Reactive manifesto	8
3.2	Concurrency	8
3.2.1	Directness and laziness	8
3.2.2	Reactive manifesto: Threads	8
3.2.3	Reactive manifesto: Futures	9
3.2.4	Reactive manifesto: Functional effect systems	9
3.2.5	Reactive manifesto: Reactive streams	9
3.3	Immutability	9
3.3.1	VAVR	10
3.4	Null-free style	10

3.5	Akka streams introduction	11
3.5.1	Source	11
3.5.2	Flow	11
3.5.3	Flow (operators)	11
3.5.4	Flow (connecting)	12
3.5.5	Sink	12
3.5.6	Materialization	12
3.5.7	Bounded processing	13
3.5.8	Custom graph stages	13
3.5.9	Use cases for reactive streams	13
3.6	Case: Kafka processing with Akka Streams	13
3.6.1	Preparation	13
3.6.2	Writing to a topic	14
3.6.3	Producer variants	14
3.6.4	Consuming from a topic	14
3.6.5	Consumer offset management	15
3.6.6	Transactions and "exactly-once" processing	15
3.7	Case: RabbitMQ processing with Akka Streams	16
3.7.1	Preparation	16
3.7.2	Writing to a topic	16
3.7.3	Producer variants	16
3.7.4	Reading from a topic	17
3.7.5	Consumer variants	17

4 Wrapping up 17

1 Introduction

1.1 About today's presentation

- *Functional JDK 17*
 - Selection of new JDK features that fit well for functional programming
 - Background and introduction to reactive streams
- A little about me, Jan Ypma
 - `jan@ypmania.net`
 - Independent software developer
 - Scala / Java, C++ (embedded), a little Rust, Lisp
- This presentation
 - <https://github.com/jypma/java-17-demo>

2 Recent new JDK features

2.1 Records (Java 16)

Declare an immutable class with constructor, getters, equals, and hashCode():

```
record Point(int x, int y) { }
```

is equivalent to

```
class Point(int x, int y) {  
    private final int x;  
    private final int y;  
  
    Point(int x, int y) {  
        this.x = x;  
        this.y = y;  
    }  
  
    int x() { return x; }  
    int y() { return y; }  
  
    @Override boolean equals(Object other) { ... }  
    @Override int hashCode() { ... }  
}
```

2.1.1 Demo

This means that we can do

```
record Point(int x, int y) {}
```

```
var p1 = new Point(1,2)
```

Use the getters:

```
"P1 is at " + p1.x() + ", " + p1.y()
```

Check for equality:

```
"P1 is equal to itself: " + p1.equals(new Point(1,2))
```

2.1.2 Summary

Use cases

- Data Transfer Objects (API client or server, database input or results)
- Nicer replacement of *tuples* (method-local records are allowed)

Limitations

- Still no immutable collections to use as field types
 - Use the **VAVR** library instead, more about that later
- No *copy* constructor (apparently being worked on)
 - Curiously, both Kotlin and Java allow this, using default arguments and a generated `copy` function:

```
val newPoint = p1.copy(x = 42)
```

2.2 Sealed classes

Interface or abstract class is marked as having a fixed set of implementations.

2.2.1 Implemented in different source files

```
package com.example.geometry;

public abstract sealed class Shape
    permits com.example.polar.Circle,
           com.example.quad.Rectangle,
           com.example.quad.simple.Square {
    // ... Shape abstract implementation...
}
```

2.2.2 Implemented in the same source file

```
abstract sealed interface Shape {
    record Circle(int diameter) implements Shape { /* ... */ }
    record Box(int width, int height) implements Shape { /* ... */ }
}
```

Each class implementing a `sealed` base can be either

- `final` (sealing the hierarchy here. All `record` classes are `final`.)
- `sealed` (implying further sub-types extending this one)
- `non-sealed` (breaking the promise of this being a sealed hierarchy)

2.2.3 Can't extend outside of what's sealed in

```
class OtherShape implements Shape {
}
```

2.3 "Pattern matching" for switch

Adds a type check to `switch` branches.

Let's recap and see how `switch` has evolved through recent Java versions.

2.3.1 Switch in Java 7

Simple replacement for `goto`, with mostly same semantics and syntax as C.

```
int value = 5;

switch(value) {
    case 1:
        System.out.println("One");
        break;
    case 5:
        System.out.println("five");
        break;
    default:
        System.out.println("Unknown");
}
```

2.3.2 Java 12: Strings, expressions

```
String day = "Tuesday";
```

We can now switch on `String`, have multiple values in one branch, and return as an expression.

Note: The following no longer compiles with Java 13+ (which requires `yield` instead of `break`).

```
switch(day) {
    case "Monday", "Tuesday":
        break "Week day";

    default:
        break "Unknown";
}
```

2.3.3 Java 12: Arrows for expressions

Instead of `break` or `yield`, arrows can be used to write a `switch` expression.

```
switch(day) {
    case "Monday", "Tuesday" -> "Week day";
    default -> "Unknown";
}
```

2.3.4 Java 13: Yield instead of break

Here's the Java 13+ equivalent:

```
switch(day) {
    case "Monday", "Tuesday":
        yield "Week day";
}
```

```

    default:
        yield "Unknown";
}

```

2.3.5 Java 17: Pattern matching objects, and guards

We can now check the type of an object, including additional constraints, right inside a `switch` case.

```
Shape shape = new Shape.Box(10, 5)
```

```

switch(shape) {
    case Shape.Circle c -> "It's a circle with diameter " + c.diameter();
    case Shape.Box b && b.width() == b.height() -> "It's a square of size " + b.width();
    case Shape.Box b -> "It's a box of size " + b.width() + " by " + b.height();
}

```

2.3.6 Exhaustiveness check

Since we defined `Shape` as a sealed class, the compiler will now inform us if we forget to add a case.

```

switch(shape) {
    case Shape.Box b && b.width() == b.height() -> "It's a square of size " + b.width();
    case Shape.Box b -> "It's a box of size " + b.width() + " by " + b.height();
}

```

2.3.7 Case branch for null

A case branch for `null` is now allowed (but, please, don't). And `default` still doesn't handle `null` (this is unchanged).

(set this to non-null to compile the `switch` below)

```
Object nothing = null;
```

```

switch (nothing) {
    case null -> "null!";
    case String s -> "String";
    default -> "Something else";
}

```

2.3.8 Limitations

- No decomposition
 - Can't match nested object graphs

```
record Drawing (Shape shape, int color)
```

```

switch (myDrawing) {
    // Does not compile:
    case Drawing(Shape.Box box, color) ->
}

```

2.4 Shenandoah GC and ZGC

2.4.1 Traditional garbage collectors

- Parallel GC
 - Stop-the-world GC for Young and Old generation
- Concurrent Mark-Sweep GC
 - Stop-the-world GC for Young, concurrent for Old generation
 - No compaction of Old generation
- G1 garbage collector
 - Stop-the-world GC for Young, concurrent mark for Old generation, stop-the-world compaction in segments
 - Configurable GC pauses: either shorter pauses, or less CPU wasted on GC
 - Default since Java 9
 - Problematic on large heaps or high allocation counts

2.4.2 ZGC and Shenandoah GC

- Scalable, low-latency GC
- No generations
- Concurrent mark *and* compaction

2.4.3 ZGC

- Since Java 11, but only on 64-bit linux (no compressed pointers)
- Store objects in ZPages (small, medium, large), compact when almost all objects in a page are dead
- Clever x86 JVM pointer tricks (*colored* pointers)
- More info on [OpenJDK wiki](#)

2.4.4 Shenandoah GC

- Developed by Red Hat
- Since Java 12 (but not in Oracle builds), but backported to 11 and 8
- Architecture independent (windows, linux and macOS)
- Derived from G1 (same marking), but divides heap into (many) *regions*
- Metadata in JVM object header
- More info on [OpenJDK wiki](#)
- So which one should I use?
 - Both ZGC and Shenandoah will probably improve your latencies
 - Try both!

2.5 macOS / AArch64 port

- Recent apple computers have 64-bit ARM processors, but don't run Linux
- There already was an `aarch64` port for Linux
- Java 17 brings native support for `aarch64` under MacOS

3 Practical reactive streams

3.1 Reactive manifesto

- **Published** in 2014, intends to push software systems to be better-behaved.
 - **Responsive:** *The system responds in a timely manner if at all possible.*
 - **Resilient:** *The system stays responsive in the face of failure.*
 - **Elastic:** *The system stays responsive under varying workload.*
 - **Message driven:** *Establish a boundary between components that ensures loose coupling, isolation and location transparency.*

3.2 Concurrency

3.2.1 Directness and laziness

- Direct value: `Person p`
 - Value is already calculated
 - This is good, we know there's no more I/O
- Direct asynchronous value: `CompletionStage<Person> p`
 - Computation already in progress: problematic
- Lazy value: `Supplier<Person> p`
 - Computation doesn't start until invoking `p.get()`
 - Nice, but not asynchronous
- Lazy asynchronous value: (no plain Java type) `"Supplier<CompletionStage<Person> p"`
 - All *Akka Streams* types are lazy and asynchronous (but multi-valued)
 - Hence, Akka can optimize and change a stream before starting it
 - * For example, adding retry behavior to stream components

3.2.2 Reactive manifesto: Threads

- Synchronous method calls
 - Hard to make *responsive* (can't really abort a thread, unless all code constantly checks time)
 - Hard to make *resilient* in Java (failure is realistically limited to exceptions, of which many are unchecked and invisible)
 - Not *message-driven* (methods return values synchronously, and/or have side effects)
- Doesn't affect *elastic*

3.2.3 Reactive manifesto: Futures

- Java calls them **CompletionStage** (`CompletionStage<T>`, `CompletableFuture<T>`,)
 - Handle to an on-going background computation
 - Hard to make *responsive* (computation already started, not cancellable in practice)
 - Even harder than Threads to make *resilient* in Java (exceptions are now hidden behind `CompletionException` plus cancellation)
 - Can model *message-driven* by having future callbacks
- Doesn't affect *elastic*

3.2.4 Reactive manifesto: Functional effect systems

- Think of "`CompletionStageRecipe<T,E>`"
 - Description of (not yet started) background computation
 - All of *responsive*, *elastic* (since description can be altered before launch) and *message-driven*
 - Very active in the Scala world (`cats-effect`, `ZIO`)
 - Not so much in plain Java or Kotlin, potentially due to missing language constructs
- Doesn't affect *elastic*

3.2.5 Reactive manifesto: Reactive streams

- Reactive streams
 - Covers a variety of independent frameworks
 - * *rxJava* (2014), porting Microsoft's "reactive extensions" to Java
 - * *Akka Streams* (2015), building on Akka with a component-based streaming framework
 - * *Project Reactor* (2015), built by Spring directly decorating `java.util.concurrent.Flow`
 - * Many others
 - Interoperability through `java.util.concurrent.Flow`
 - * Low-level
 - We'll look at Akka Streams today
 - * Trivially *responsive* (real time is a core element of streams)
 - * *Resilient* due to well-defined error propagation and handling
 - * Gives some *elastic* guarantees due to bounded processing (more on that later)
 - * Integrates well in *message-driven* architectures (native actors support)

3.3 Immutability

- Asynchronous processing on data needs guarantees
 - Locks? Not if each and every data object is processed concurrently.
- "*I promise I won't change this object anymore*" just isn't cutting it
- Need actual immutability
 - Have compiler help guaranteeing objects won't be changed

- No setters
- `record` anyone?
- Can't use `java.util.List` or `java.util.Map`

3.3.1 VAVR

- **Functional library** for Java, focusing on immutable values
- **JavaDoc** shows collection, control and concurrency primitives

Create an immutable sequence:

```
Seq<Integer> seq = Vector.of(1, 2, 3)
```

```
seq.forEach(i -> System.out.println(i))
```

- All VAVR collections are *persistent data structures*, for example
 - `List` (single-linked list)
 - `Vector` (bit-mapped trie)
 - `HashMap` (hash array mapped trie)

3.4 Null-free style

- Nobody likes `NullPointerException`
- Reactive streams, and most functional libraries, don't allow (or like) `null` as values
- So, why are we still using `null` to indicate optionality?
 - Use `java.util.Optional` or the more powerful `io.vavr.control.Option` (or `io.vavr.control.Either`) instead.

```
Option<User> getUserIfExists(userId: long) {
  // ...
}
```

- In case of optional method arguments, consider method overloading instead of passing `null` (but `Option` is also fine here).

```
void saveUser(String userName, String petName) {
  // Save a user who signed up together with their pet.
}
```

```
void saveUser(String userName) {
  // Save a user who signed up by themselves.
}
```

- In short
 - The word `null` should never occur in your pull requests for new code
 - Only exception is interacting with external `null`-loving libraries

3.5 Akka streams introduction

- **Akka Streams:** Composable reactive streams framework
- Implemented on top of Akka *actors* (but invisibly so). You need an `ActorSystem` to launch streams:

```
ActorSystem system = ActorSystem.create("Demo")
```

- Streams form a graph, built using components called *graph stages*
 - Type-safe *input(s)* and/or *output(s)*
 - Number of inputs and outputs defines its *shape*
- Stream objects are descriptions only, and need to be *materialized* to actually do something

3.5.1 Source

[source.gif](#)

- Has a single output of type T, no inputs
- Emits elements

For example, a source that emits the same element every second:

```
Source<String, Cancellable> everySecond = Source.tick(Duration.ofSeconds(1),  
↳ Duration.ofSeconds(1), "tick!")
```

Or a source that emits all integers up to one million, as fast as the stream can use them:

```
Source<Integer, NotUsed> integers = Source.range(1, 1000000)
```

3.5.2 Flow

[flow.gif](#)

- Has a single input of type T, and one output of type U
- Typically emits elements on its output as it receives them in the input

For example, a flow that converts integers to strings:

```
Flow<Integer, String, NotUsed> intToString = Flow.<Integer>create().  
map(i -> i.toString())
```

3.5.3 Flow (operators)

But we have more complex, useful operators. For example, process a sliding window of 10 elements: (*we'll map to VAVR's `Vector` to ensure immutability*)

```
Flow<Integer, Seq<Integer>, NotUsed> intSliding = Flow.<Integer>create().  
sliding(1, 10).  
map(Vector::ofAll)
```

Or, group elements up to a certain count, *OR* until some time has elapsed:

```
Flow<Integer, Seq<Integer>, NotUsed> intGrouped = Flow.<Integer>create().  
groupedWithin(256, Duration.ofSeconds(1)).  
map(Vector::ofAll)
```

3.5.4 Flow (connecting)

- Connecting a Flow to a Source (of compatible type) can be viewed as a Source (of the Flow's output type)

For example, let's hook up our `integers` source to the `intToString` flow:

```
Source<String,NotUsed> strings = integers.via(intToString)
```

In order to test, let's print the first 10 elements which that flow produces.

```
strings.  
  take(10).  
  runForeach(System.out::println, system).  
  toCompletableFuture().get(1, TimeUnit.SECONDS)
```

3.5.5 Sink

[sink.gif](#)

- Has a single input of type T
- Typically "consumes" the elements

```
Sink<String, CompletionStage<Done>> printStrings = Sink.<String>foreach(s ->  
  ↪ System.out.println(s))
```

- Connecting a Source to a Sink leaves no inputs or outputs
 - Akka calls this a `RunnableGraph`

```
RunnableGraph<NotUsed> graph = strings.to(printStrings)
```

- We won't run the above graph, since there's no `CompletionStage` indicating when it's done (only `NotUsed`)

3.5.6 Materialization

- Instances of graphs (`Source`, `Sink`, ...) are *descriptions*, and don't run yet
- Need to invoke `RunnableGraph.run()` (or one of the shorthands on `Source`) to actually start a stream
- Running a stream gives a *materialized value*
 - `Source<T, M>`. emits elements of type T, results in a value M when started
 - `Sink<T, M>`. consumes elements of type T, results in a value M when started
 - `RunnableGraph<M>`. results in a value M when started (`.run()` returns M)
- Now, we can construct `graph` again, but this time use the materialized value of the `sink`
 - By default, `.to()` uses the materialized value of the `source`

```
RunnableGraph<CompletionStage<Done>> graph = strings.take(10).toMat(printStrings,  
  ↪ (sourceMat, sinkMat) -> sinkMat)
```

```
graph.run(system).toCompletableFuture().get(1, TimeUnit.SECONDS)
```

3.5.7 Bounded processing

- When writing data processing software, always make sure to be explicit in how much *of each* you want in memory
- Akka makes this explicit wherever possible
 - `groupedWithin` takes a maximum amount of elements AND a duration. There is no variant that only takes a duration.

```
source.groupedWithin(100, Duration.ofSeconds(1))
```

- `groupBy(Integer maxStreams, Function<T,K> key)` (grouping substreams by key) needs to specify the maximum number of open streams
- `mapAsync` (allowing to map each element to a `CompletionStage`'s result) needs to specify the number of in-flight elements
- Akka helps you towards bounded processing

3.5.8 Custom graph stages

- Writing your own `Source`, `Flow` or `Sink` is easy and well-documented
- These are ideal building blocks for data-processing systems
 - Encapsulate resource handling inside your building block
 - Well-defined error handling and propagation

3.5.9 Use cases for reactive streams

Good reasons to reach for reactive streams:

- Variance in iteration size
 - Being able to handle, simultaneously, both *many small* requests but also *few large* requests with the same code
- Heterogeneous systems
- Predictable memory usage

3.6 Case: Kafka processing with Akka Streams

3.6.1 Preparation

- Kafka is running locally, started from `docker-compose.yml`
- Let's make sure we have an empty topic to play with:

```
kafkactl delete topic demo 2>/dev/null
kafkactl create topic demo
```

- Akka can make use of Kafka through the `Alpakka Kafka` library

3.6.2 Writing to a topic

- Let's use akka's `Producer.plainSink` in a simple example

```
void writeToTopic() throws Exception {
    final ProducerSettings<String, String> producerSettings =
        ProducerSettings.create(system, new StringSerializer(), new StringSerializer())
            .withBootstrapServers("localhost:9092");

    Source.range(1, 10)
        .map(number -> number.toString())
        .map(value -> new ProducerRecord<String, String>("demo", value))
        .runWith(Producer.plainSink(producerSettings), system)
        .toCompletableFuture().get(10, TimeUnit.SECONDS);
}

writeToTopic()
```

- Let's see if they arrived:

```
kafkactl consume demo --from-beginning --exit
```

3.6.3 Producer variants

The Alpakka Producer class has *several ways* of defining a Kafka producer.

- `Producer.plainSink`: Sends `ProducerMessage` objects to Kafka
 - Suitable when sending to Kafka is the last step in a stream
- `Producer.flexiFlow`: Sends `Envelope` to Kafka, and passes it on down-stream
 - An `Envelope` can potentially contain more than one Kafka message, and an arbitrary *context* object
 - Useful when you need to do more after sending to Kafka
- `Producer.committableSink`: Automatically *commits* messages read from another Kafka topic
 - Useful in *consume - process - produce* type flows

3.6.4 Consuming from a topic

- Let's use the Alpakka `Consumer.plainSource` in a simple example

```
Seq<String> readFromTopic() throws Exception {
    final ConsumerSettings<String, String> consumerSettings =
        ConsumerSettings.create(system, new StringDeserializer(), new
            ↪ StringDeserializer())
            .withBootstrapServers("localhost:9092")
            .withGroupId("group1")
            .withProperty(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

    return Consumer.plainSource(consumerSettings, Subscriptions.topics("demo"))
        .take(1)
}
```

```

        .map(record -> record.value())
        .runWith(Sink.seq(), system)
        .thenApply(Vector::ofAll)
        .toCompletableFuture()
        .get(20, TimeUnit.SECONDS);
    }

```

(demo is unfortunately not working due to JShell limitations)

readFromTopic()

3.6.5 Consumer offset management

Kafka can store the offset for consumer groups, or consumers can provide (and store) it themselves.

- Store offset in Kafka
 - `Consumer.committableSource`
- No offset management
 - `Consumer.plainSource(settings, Subscriptions.topics("topic"))`
- Do your own offset management
 - `Consumer.plainSource(settings, Subscriptions.assignmentWithOffset(new TopicPartition("topic", partition0), fromOffset))`
 - After each element, store its partition and offset in your own storage

3.6.6 Transactions and "exactly-once" processing

- Recent Kafka versions implement an extension that allows clients to atomically
 - Consume from one topic
 - Produce results to another topic
- Kafka refers to this both as *transactions* and *exactly-one processing*
- This feature can be used from Akka using the `Alpakka Transactional` class, e.g.

```

Transactional.source(consumerSettings, Subscriptions.topics(sourceTopic))
    .via(business())
    .map(
        msg ->
        ProducerMessage.single(
            new ProducerRecord<>(targetTopic, msg.record().key(), msg.record().value()),
            msg.partitionOffset())
    ).toMat(
        Transactional.sink(producerSettings, transactionalId),
        Consumer::createDrainingControl)
    .run(system);

```

- `PartitionOffset` holds the partition number and offset of the originally consumed message
- This is passed as *context* argument to the `ProducerRecord`

3.7 Case: RabbitMQ processing with Akka Streams

3.7.1 Preparation

- RabbitMQ is running locally, started from [docker-compose.yml](#)
- Communication is over AMQP, using akka's [Alpakka AMQP](#) library

3.7.2 Writing to a topic

```
Seq<WriteResult> writeToTopic() throws Exception {  
    var settings = AmqpWriteSettings.create(AmqpLocalConnectionProvider.getInstance())  
        .withRoutingKey("demo-queue")  
        .withDeclaration(QueueDeclaration.create("demo-queue"))  
        .withBufferSize(10)  
        .withConfirmationTimeout(Duration.ofMillis(200));  
  
    return Source.range(1, 10)  
        .map(number -> number.toString())  
        .map(value -> WriteMessage.create(ByteString.fromString(value)))  
        .via(AmqpFlow.createWithConfirm(settings))  
        .runWith(Sink.seq(), system)  
        .thenApply(Vector::ofAll)  
        .toCompletableFuture().get(10, TimeUnit.SECONDS);  
}
```

writeToTopic()

3.7.3 Producer variants

RabbitMQ (and its underlying AMQP protocol) allows varying degrees of consistency when producing messages.

- *Fire-and-forget* : Fastest performance, but messages may be lost in case of broker or network issues
 - Use `AmqpFlow.apply`
- *Publisher confirms*: Asynchronous message from RabbitMQ to client (after fsync)
 - Use `AmqpFlow.withConfirm` (setting `bufferSize` to the allowed number of parallel in-flight messages)
 - Use `AmqpFlow.withConfirmUnordered` for maximum throughput, sacrificing ordering guarantees
- *Transactions*
 - Traditionally considered "slow" by RabbitMQ
 - Not directly supported by the Alpakka library (just use publisher confirms)

3.7.4 Reading from a topic

```
Seq<String> readFromTopic() throws Exception {
    var bufferSize = 10;
    Source<ReadResult, NotUsed> amqpSource =
        AmqpSource.atMostOnceSource(
            NamedQueueSourceSettings.create(AmqpLocalConnectionProvider.getInstance(),
                ↪ "demo-queue")
                .withDeclaration(QueueDeclaration.create("demo-queue"))
                .withAckRequired(false),
            bufferSize);

    return amqpSource.take(10)
        .map(readResult -> readResult.bytes().utf8String())
        .runWith(Sink.seq(), system)
        .thenApply(Vector::ofAll)
        .toCompletableFuture()
        .get(1, TimeUnit.SECONDS);
}

readFromTopic()
```

3.7.5 Consumer variants

RabbitMQ (and its underlying AMQP protocol) allows varying degrees of consistency when consuming messages.

- *Consumer acknowledgement*
 - Consumers send an `ack` message to RabbitMQ to indicate that they've successfully processed a message
 - Consumers can `ack` all messages up to the current one with one confirmation
 - Use `AmqpSource.committableSource`, process each element, and then invoke `.ack()` on it
 - * `.mapAsync(committableReadResult -> committableReadResult.ack())`
- *Automatic acknowledgement*
 - Akka can automatically acknowledge messages as soon as they're read
 - Use `AmqpSource.atMostOnceSource`
- *Transactions*
 - Traditionally considered "slow" by RabbitMQ
 - Not directly supported by the Alpakka library (just use consumer acknowledgement)

4 Wrapping up

- Thanks for your participation!
- Any final thoughts / questions?
- Curious how this presentation was made?
 - Attend my talk at [EmacsConf 2021](#)